

coursework_01

February 2, 2026

1 Coursework 1: Image filtering

The coursework includes coding questions and/or written questions. Please read both the text and the code in this notebook to get an idea what you are expected to implement.

1.1 What is expected?

- Complete and run the code using `jupyter-lab`.
- Export (File | Save and Export Notebook As...) the notebook as a PDF file, which contains your code, results and answers, and upload the PDF file onto [Scientia](#).
- Instead of clicking the Export button, you can also run the following command instead:
`jupyter nbconvert coursework_01.ipynb --to pdf`
- If Jupyter complains issues during exporting, it is likely that `pandoc` or `latex` is not installed, or their paths have not been included. You can install the relevant libraries and retry. Alternatively, use the Print function of your browser to export the PDF file.
- If Jupyter-lab does not work for you at the end, alternatively, you can use Google Colab to write the code and export the PDF file.

1.2 Dependencies:

You may need to install [Jupyter-Lab](#) and other libraries used in this coursework, such as by running the command: `pip3 install [package_name]`

```
[44]: # Import libraries (provided)
import imageio.v3 as imageio
import numpy as np
import matplotlib.pyplot as plt
import scipy
import scipy.signal
import math
import time
```

1.3 Q1. Moving average filtering (20 points).

Read the provided clean image, add noise to the image and design a moving average filter for denoising.

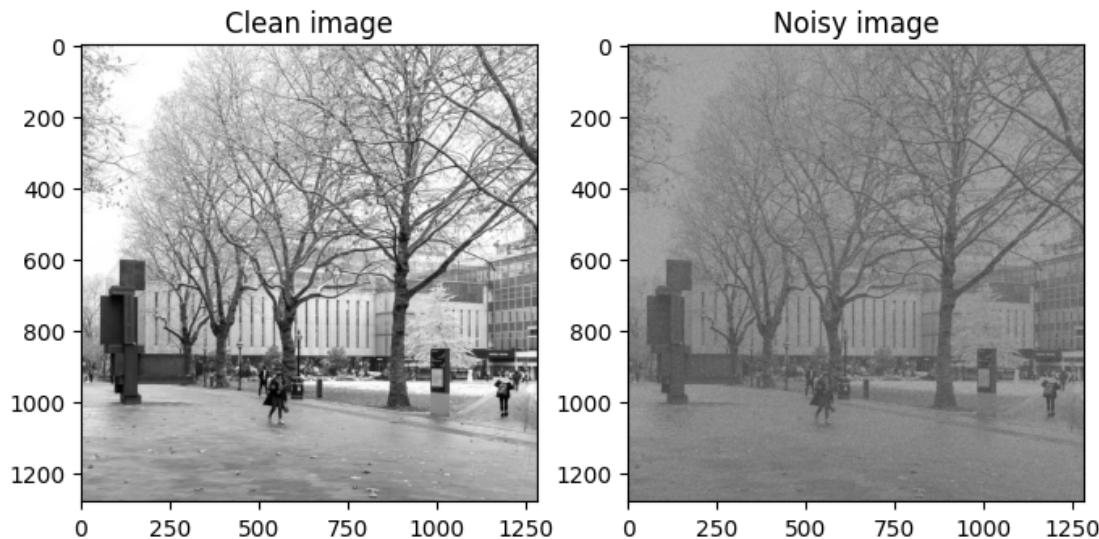
You are expected to design the kernel of the filter and then perform 2D image filtering using the function `scipy.signal.convolve2d()`.

```
[45]: # Read the image (provided)
image = imageio.imread('campus_snow.jpg')

# Corrupt the image with Gaussian noise (provided)
noise_mu = 0
noise_sigma = 50
noise = np.random.normal(noise_mu, noise_sigma, image.shape)
image_noisy = image + noise

# Visualise the images (provided)
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Clean image')

plt.subplot(1, 2, 2)
plt.imshow(image_noisy, cmap='gray')
plt.title('Noisy image')
plt.gcf().set_size_inches(8, 4)
```



1.3.1 Q1.1 Filter the noisy image using a 5x5 moving average filter. Display the filtered image.

```
[46]: # Design the filter h
### Insert your code ####
h = np.array([[1/25, 1/25, 1/25, 1/25, 1/25] for i in range(5)])
```

```

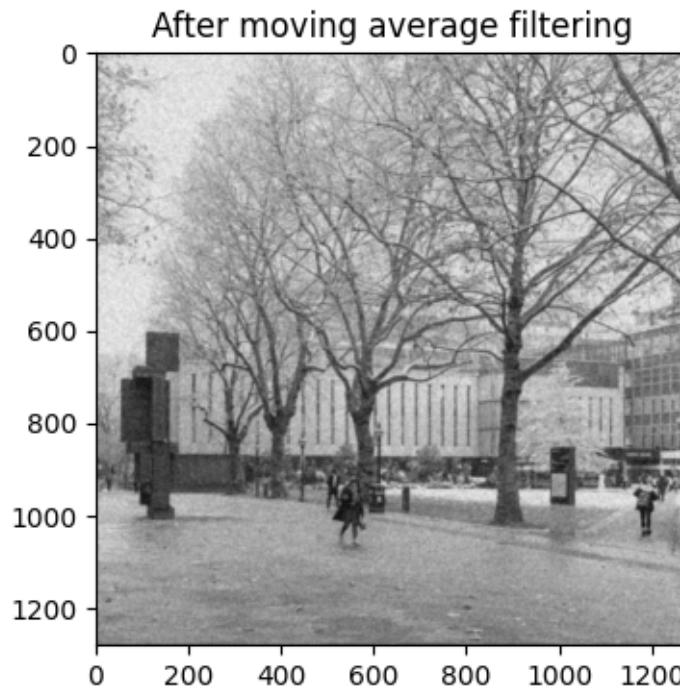
# Convolve the noisy image with h using scipy.signal.convolve2d function
### Insert your code ####
image_filtered = scipy.signal.convolve2d(image_noisy, h, 'same')

# Print the filter (provided)
print('Filter h = {0}'.format(h))

# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('After moving average filtering')
plt.gcf().set_size_inches(4, 4)

```

Filter h = [[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]]



1.3.2 Q1.2 Assess the quality of the filtered image using a quantitative metric, the peak signal-to-noise ratio (PSNR).

For this case, the pixel intensity of the image is represented using the uint8 format, with the peak value to be 255. The PSNR is defined as,

$$\text{PSNR} = 10 \cdot \log_{10} \frac{255^2}{\frac{1}{N} \sum_x [J(x) - I(x)]^2}$$

where x denotes the pixel index, N denotes the total number of pixels in the image, J denotes the filtered i.e. denoised image and I denotes the ground truth clean image. The denominator of the term within the logarithm operator is the mean squared error between I and J .

You can find more detail about PSNR [here](#).

```
[47]: # Implement the PSNR function
def eval_psnr(I, J):
    # I: the ground truth clean image (peak value: 255 for uint8 data format)
    # J: the denoised image
    #
    # return: the PSNR metric (unit: dB)
    N = np.size(I)

    Sum = np.sum(np.square(J - I))

    ### Insert your code ####
    psnr = 10 * np.log10(np.square(255) / ((1/N) * Sum))
    return psnr

# Evaluate the PSNR for the filtered image (provided)
psnr = eval_psnr(image, image_filtered)

# Print the PSNR (provided)
print('PSNR = {:.2f} dB'.format(psnr))
```

PSNR = 18.39 dB

1.4 Q2. Gaussian filtering (70 points).

1.4.1 Q2.1 Implement a function that constructs a 2D Gaussian filter given the parameter σ .

```
[48]: # Implement the Gaussian filter
def gaussian_filter_2d(sigma):
    # sigma: the parameter sigma for the Gaussian kernel (unit: pixel)
    #
    # return: a 2D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))
    sz = 2 * rad + 1

    constant = 1/(2*np.pi*np.square(sigma))

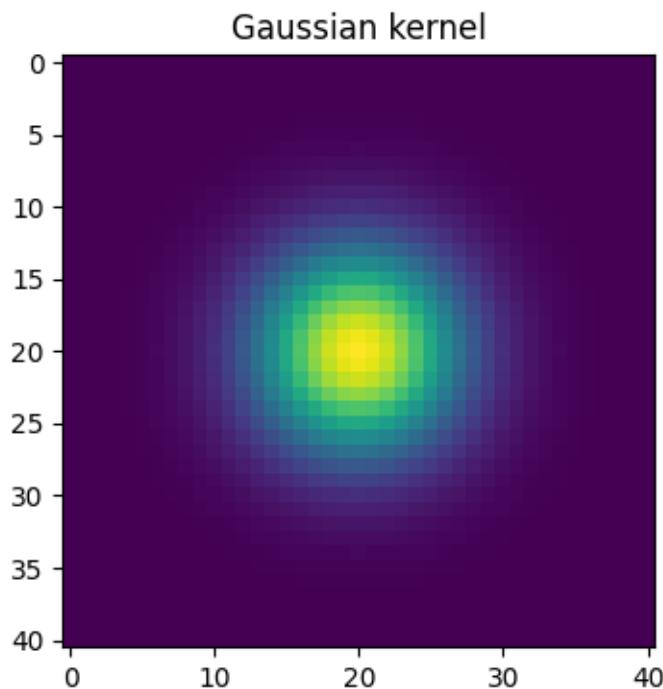
    # Calculate the filter weights
    ### Insert your code ####
```

```

h = np.array([[constant * (np.e ** -((np.square(i - (rad)) + np.square(j - (rad)))/ (2 * np.square(sigma)))) for i in range(sz)] for j in range(sz)])
return h

# Visualise the Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)
plt.imshow(h)
plt.title('Gaussian kernel')
plt.gcf().set_size_inches(4, 4)

```



1.4.2 Q2.2 Perform Gaussian filtering ($\sigma = 5$ pixels) for the noisy image, evaluate the computational time for Gaussian filtering and display the filtered image.

```
[49]: # Construct the Gaussian filter (provided)
sigma = 5
h = gaussian_filter_2d(sigma)

# Perform Gaussian filtering and count time
### Insert your code ####
start = time.time()
image_smoothed_2d = scipy.signal.convolve2d(image_noisy, h, 'same')
end = time.time()
```

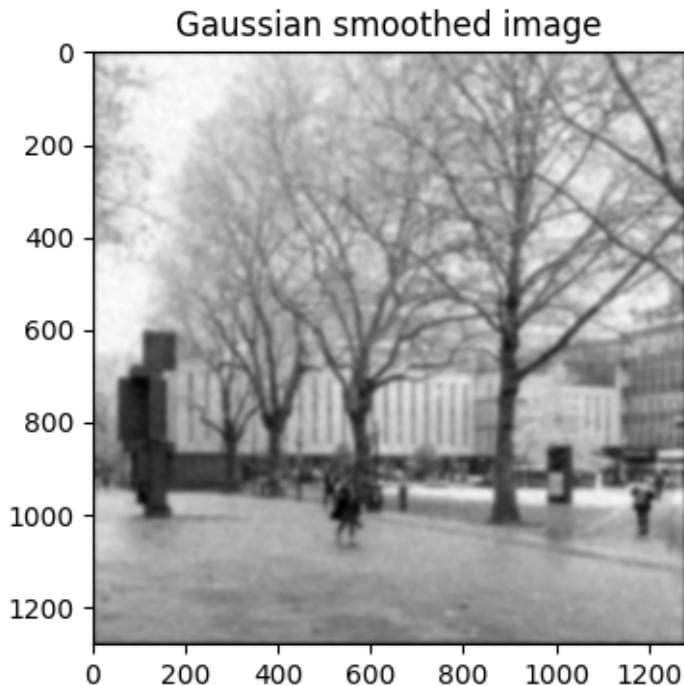
```

print(end - start)

# Visualise the filtered image (provided)
plt.imshow(image_smoothed_2d, cmap='gray')
plt.title('Gaussian smoothed image')
plt.gcf().set_size_inches(4, 4)

```

2.398364305496216



1.4.3 Q2.3 Implement a function that generates a 1D Gaussian filter given the parameter σ . Construct 1D Gaussian filters along x-axis and y-axis respectively.

```

[50]: # Implement the 1D Gaussian filter
def gaussian_filter_1d(sigma):
    # sigma: the parameter sigma in the Gaussian kernel (unit: pixel)
    #
    # return: a 1D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))
    sz = 2 * rad + 1

    constant = 1/(np.sqrt(2*np.pi)*sigma)

```

```

# Calculate the filter weights
### Insert your code ####
h = [[constant * (np.e ** -(np.square(i - (rad)) / (2 * np.square(sigma))))]
    ↪for i in range(sz)]]
return h

# sigma = 5 pixel (provided)
sigma = 5

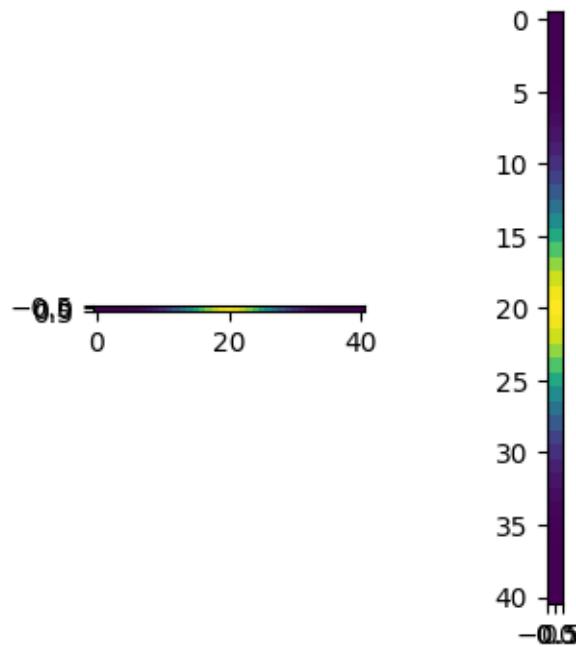
# Construct the Gaussian filter along x-axis. Its shape is (1, sz).
### Insert your code ####
h_x = gaussian_filter_1d(sigma)

# Construct the Gaussian filter along y-axis. Its shape is (sz, 1).
### Insert your code ####
h_y = np.transpose(gaussian_filter_1d(sigma))

# Visualise the filters (provided)
plt.subplot(1, 2, 1)
plt.imshow(h_x)

plt.subplot(1, 2, 2)
plt.imshow(h_y)
plt.gcf().set_size_inches(4, 4)

```



1.4.4 Q2.4 Perform Gaussian filtering ($\sigma = 5$ pixels) using two separable filters and evaluate the computational time for separable Gaussian filtering. Compare the smoothed image using separable filtering to the smoothed image using a single 2D Gaussian filter.

```
[51]: # Perform separable Gaussian smoothing and count time
### Insert your code ####
start = time.time()
image_smoothed = scipy.signal.convolve2d(image_noisy, h_x, 'same')
image_smoothed = scipy.signal.convolve2d(image_smoothed, h_y, 'same')
end = time.time()

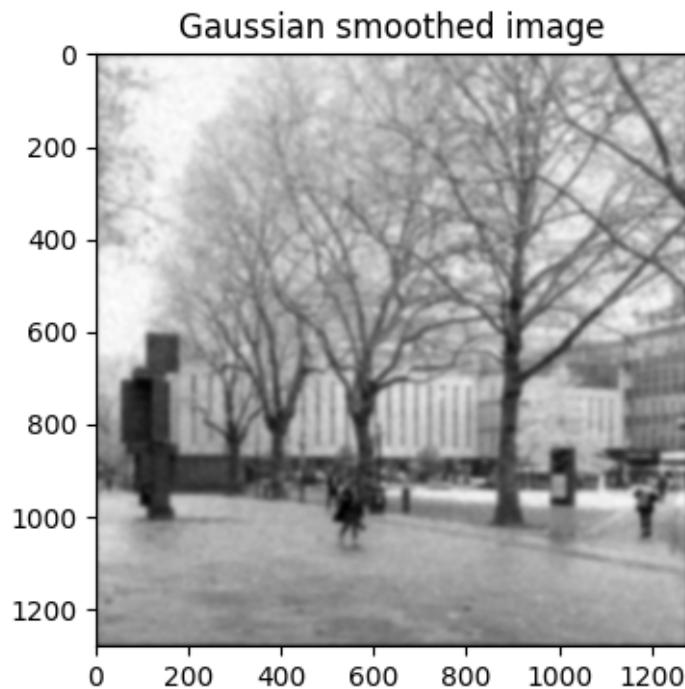
print(end - start)

# Report the difference between the separably filtered image and the image
# ↪filtered by a single 2D Gaussian filter (provided)
diff = image_smoothed - image_smoothed_2d
print('Mean absolute difference = {:.6f}'.format(np.mean(np.abs(diff))))
```



```
plt.imshow(image_smoothed, cmap='gray')
plt.title('Gaussian smoothed image')
plt.gcf().set_size_inches(4, 4)
```

```
0.3275880813598633
Mean absolute difference = 0.000000
```



1.4.5 Q2.5 Perform Gaussian smoothing for the same noisy image, assess the quality of the Gaussian smoothed image using PSNR, when different sigma values are used.

```
[52]: # A list of sigma values (provided)
list_sigma = np.arange(0.5, 5, 0.5)

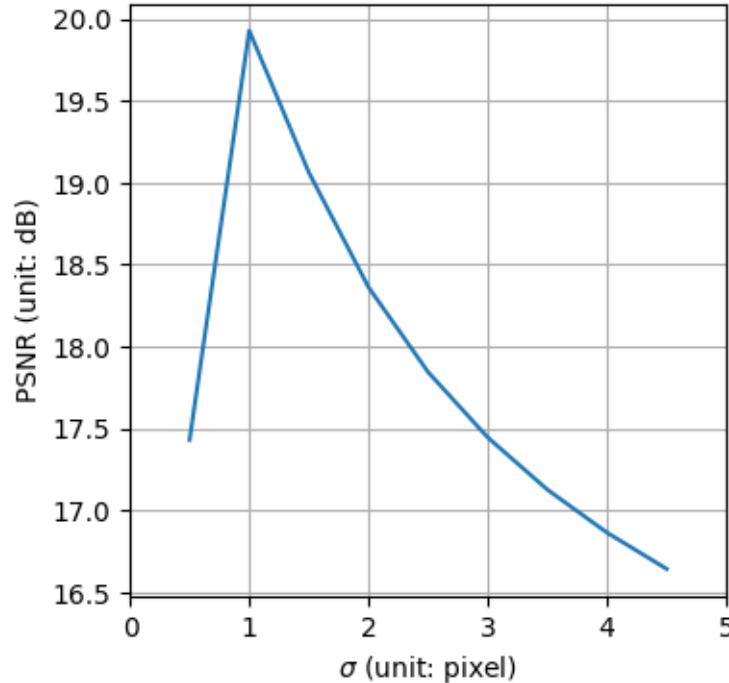
# Perform Gaussian smoothing with different sigma values and record the PSNR
# values
### Insert your code ####
list_psnr = []
for sigma in list_sigma:
    h_x = gaussian_filter_1d(sigma)
    h_y = np.transpose(gaussian_filter_1d(sigma))

    image_smoothed = scipy.signal.convolve2d(image_noisy, h_x, 'same')
    image_smoothed = scipy.signal.convolve2d(image_smoothed, h_y, 'same')

    list_psnr.append(eval_psnr(image, image_smoothed))

# Plot the PSNR metric against sigma (provided)
plt.plot(list_sigma, list_psnr)
plt.xlim([0, 5])
plt.xlabel('$\sigma$ (unit: pixel)')
plt.ylabel('PSNR (unit: dB)')
plt.grid()
plt.gcf().set_size_inches(4, 4)
```

```
<>:20: SyntaxWarning: "\s" is an invalid escape sequence. Such sequences will
not work in the future. Did you mean "\\\s"? A raw string is also an option.
<>:20: SyntaxWarning: "\s" is an invalid escape sequence. Such sequences will
not work in the future. Did you mean "\\\s"? A raw string is also an option.
/var/folders/g1/43z3gnl916gf4ljxlypcfn280000gn/T/ipykernel_3989/1241216230.py:20
: SyntaxWarning: "\s" is an invalid escape sequence. Such sequences will not
work in the future. Did you mean "\\\s"? A raw string is also an option.
    plt.xlabel('$\sigma$ (unit: pixel)')
```



1.4.6 Q2.6 Implement 3x3 Sobel filters, perform Sobel filtering for the noisy image, and display the gradient magnitude map.

```
[53]: # Construct the Sobel filters
### Insert your code ####
sobel_x = np.array([[1, 0, -1], [2, 0, -2], [1, 0, -1]])
sobel_y = np.array([[1, 2, 1], [0, 0, 0], [-1, -2, -1]])

# Print the filters (provided)
print('Sobel_x = {}'.format(sobel_x))
print('Sobel_y = {}'.format(sobel_y))

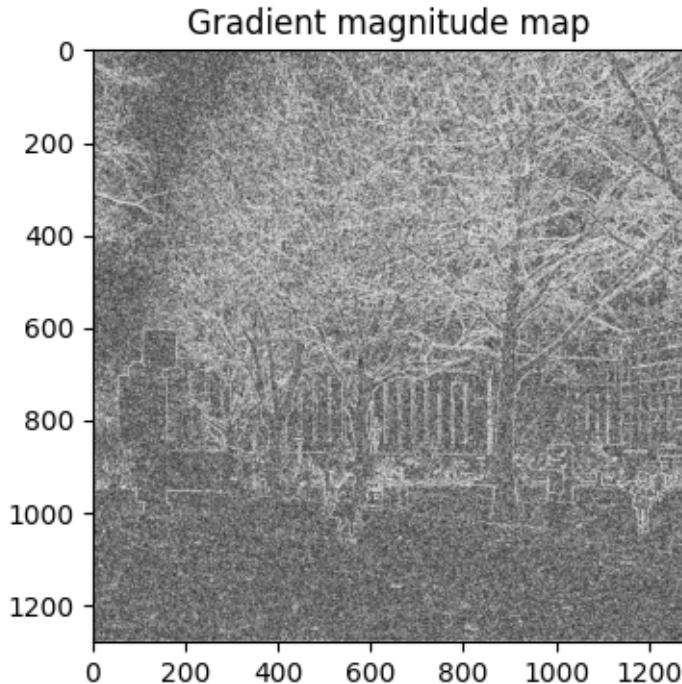
# Sobel filtering for the noisy image
### Insert your code ####
x_derivatives = scipy.signal.convolve2d(image_noisy, sobel_x, 'same')
y_derivatives = scipy.signal.convolve2d(image_noisy, sobel_y, 'same')

# Calculate the gradient magnitude
### Insert your code ####
grad_mag_noisy = np.sqrt(np.square(x_derivatives) + np.square(y_derivatives))

# Display the magnitude map (provided)
plt.imshow(grad_mag_noisy, cmap='gray', vmin=0, vmax=500)
plt.title('Gradient magnitude map')
```

```
plt.gcf().set_size_inches(4, 4)
```

```
Sobel_x = [[ 1  0 -1]
           [ 2  0 -2]
           [ 1  0 -1]]
Sobel_y = [[ 1  2  1]
           [ 0  0  0]
           [-1 -2 -1]]
```



1.4.7 Q2.7 Perform Gaussian smoothing for the noisy image, followed by Sobel filtering and display the gradient magnitude map.

```
[54]: # Parameter for the Gaussian filter (provided)
sigma = 5

# Gaussian smoothing
### Insert your code ####
h_x = gaussian_filter_1d(sigma)
h_y = np.transpose(gaussian_filter_1d(sigma))

image_smoothed = scipy.signal.convolve2d(image_noisy, h_x, 'same')
image_smoothed = scipy.signal.convolve2d(image_smoothed, h_y, 'same')

# Sobel filtering
```

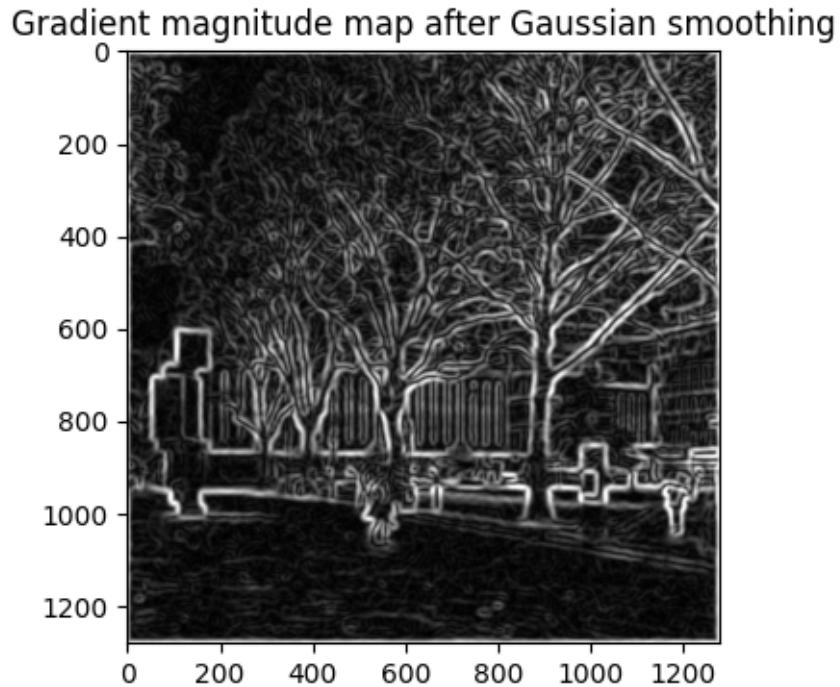
```

### Insert your code ####
x_derivatives = scipy.signal.convolve2d(image_smoothed, sobel_x, 'same')
y_derivatives = scipy.signal.convolve2d(image_smoothed, sobel_y, 'same')

# Calculate the gradient magnitude
### Insert your code ####
grad_mag = np.sqrt(np.square(x_derivatives) + np.square(y_derivatives))

# Display the magnitude map (provided)
plt.imshow(grad_mag, cmap='gray', vmin=0, vmax=100)
plt.title('Gradient magnitude map after Gaussian smoothing')
plt.gcf().set_size_inches(4, 4)

```



1.5 Q3. Implement image filters using Pytorch (10 points).

Pytorch is a machine learning framework that supports filtering and convolution.

The **Conv2D** operator takes an input array of dimension $N \times C_1 \times X \times Y$, applies the filter and outputs an array of dimension $N \times C_2 \times X \times Y$. Here, since we only have one image with one colour channel, we will set $N=1$, $C_1=1$ and $C_2=1$. You can read the documentation of Conv2D for more detail.

```
[55]: # Import libraries (provided)
import torch
```

1.5.1 Q3.1 Expand the dimension of the noisy image into 1x1xXxY and convert it to a Pytorch tensor.

```
[56]: # Expand the dimension of the numpy array
#### Insert your code ####
torch_image_noisy = np.array([[image_noisy]]).astype(np.float32)

# Convert to a Pytorch tensor using torch.from_numpy
#### Insert your code ####
tensor_noisy = torch.from_numpy(torch_image_noisy)
```

1.5.2 Q3.2 Create a Pytorch Conv2D filter, set its kernel to be a 2D Gaussian filter, perform filtering, report computational time and display the result.

```
[59]: # A 2D Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)
h = np.array([[[h]]]).astype(np.float32)

# Construct the Conv2D filter
#### Insert your code ####
kernel = torch.from_numpy(h)
conv = torch.nn.Conv2d(1,1, kernel_size=h.shape, stride=1, padding=h.shape[-1]//  
    ↵2, bias=False)
kernel.requires_grad = True

with torch.no_grad():
    conv.weight = torch.nn.Parameter(kernel)

# Filtering and assess computational time
#### Insert your code ####

start = time.time()
image_filtered = conv(tensor_noisy)
end = time.time()

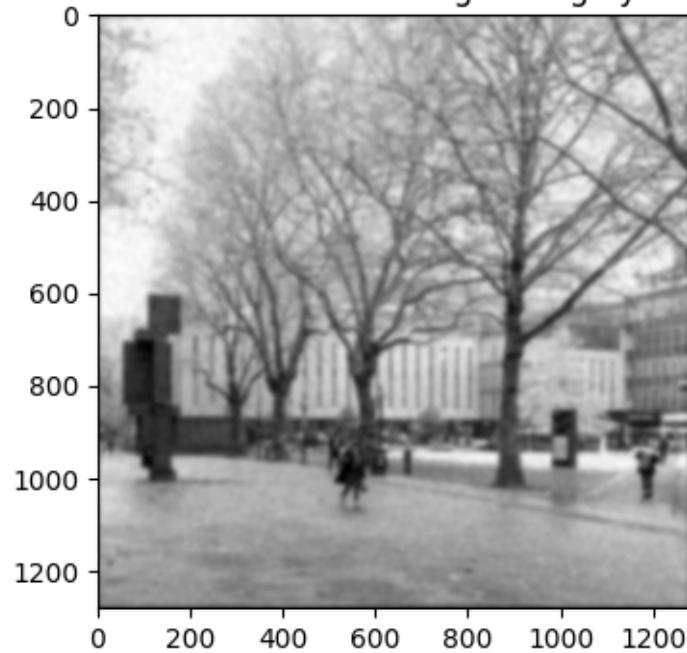
print(end - start)

image_filtered = image_filtered.view(image_filtered.shape[-2:])

# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('Gaussian smoothed image using PyTorch')
plt.gcf().set_size_inches(4, 4)
```

1.481459140777588

Gaussian smoothed image using PyTorch



[]: